

PRACTICE PROBLEMS

OPERATING SYSTEMS: INTERNALS AND DESIGN PRINCIPLES EIGHTH EDITION

WILLIAM STALLINGS

Copyright 2014: William Stallings

ACKNOWLEDGEMENTS

I would like to thank the following people for providing problems for this collection: Prof. Tracy Camp (Colorado School of Mines); Prof. Oskars J. Rieksts (Kutztown University); Dr. I-Ling Ye (U. of Texas, Dallas); Prof. Thomas N. Scanlan (U. of Wisconsin, Platteville); Klara Nahrstedt (U. of Illinois, Urbana-Champaign); Prof. Douglas Niehaus (U. of Kansas).

TABLE OF CONTENTS

Chapter 1	Computer System Overview.....	4
Chapter 2	Operating System Overview.....	6
Chapter 3	Process Description and Control	7
Chapter 5	Concurrency: Mutual Exclusion and Synchronization...	8
Chapter 6	Concurrency: Deadlock and Starvation	12
Chapter 7	Memory Management.....	15
Chapter 8	Virtual Memory	18
Chapter 9	Uniprocessor Scheduling.....	25
Chapter 11	I/O Management and Disk Scheduling	27
Chapter 12	File Management	29

CHAPTER 1 COMPUTER SYSTEM OVERVIEW

1.1 Discuss the mechanism for interrupt handling of I/O operations. Be sure to differentiate between hardware (or firmware) functions and software (OS) functions.

1.2 Consider the following pseudo assembly code for computing $c = a + b$. Assume that a , b , and c are assigned to consecutive memory “words” (memory is generally addressed byte by byte and assume that a word is 4 bytes) and address for “ a ” is $0x0000ec00$. Also, we have $a = 22$, $b = 158$, and $c = 0$ at the starting time. Assume that the first instruction of the code is stored in $0x0000b128$. Also, each instruction has the opcode in the first byte (most significant byte) and the remaining 3 bytes specify the corresponding address. The opcode for store is 1, load is 2, and add is 3.

```
0x0000b128 load a
0x0000b12c add b
0x0000b130 store c
```

a. Show the memory addresses and contents for all the instructions and data involved. Use the format as follows to express your answer (but the following is not the answer). For all data, use hexadecimal representation.

addresses	contents
0x00002104	0x00000001
0x00002108	0x00000002

.....

b. Write the micro instructions for the code segment. Assume that current PC (program counter) is $0x00001018$. For each micro-instruction, also indicate the data that is transferred (if it is a memory access). For all data, use the hexadecimal representation. The following are the first two micro-instructions and the data transferred. Complete the rest.

Micro-instructions	data
PC \rightarrow MAR	$0x0000b128$
M \rightarrow MBR	$0x0200ec00$

1.3 A computer has a cache, main memory, and a disk used for virtual memory. An access to the cache takes 10 ns. An access to main memory takes 100 ns. An access to the disk takes 10,000 ns. Suppose the cache hit ratio is 0.9 and the main memory hit ratio is 0.8. What is

the effective access time (EAT) in ns required to access a referenced word on this system?

- 1.4 a.** Consider a main memory system that consists of a number of memory modules attached to the system bus, which is one word wide. When a write request is made, the bus is occupied for 100 nanoseconds (ns) by the data, address, and control signals. During the same 100 ns, and for 500 ns thereafter, the addressed memory module executes one cycle accepting and storing the data. The (internal) operation of different memory modules may overlap in time, but only one request can be on the bus at any time. What is the maximum number of stores that can be initiated in one second?
- b.** Sketch a graph of the maximum write rate (words per second) as a function of the module cycle time, assuming eight memory modules and a fixed bus busy time of 100 ns.

CHAPTER 2 OPERATING SYSTEM OVERVIEW

2.1 Suppose a short-term scheduling algorithm favors those processes that have used little processor time in the recent past.

- a. Explain why this algorithm favors I/O-bound processes.
- b. Explain why this algorithm does not permanently deny processor time to CPU-bound processes.

2.2 The classical batch processing system ignores the cost of increased waiting time for users. Consider a single batch characterized by the following parameters:

- M average mounting time
- T average service time per job
- N number of jobs
- S unit price of service time
- W unit price of waiting time per user

- a. Show that the optimal batch size that minimizes the cost of service time and waiting time per user (within a single batch) is

$$N_{opt} = \sqrt{\frac{M}{T} \frac{S}{W}}$$

- b. In an installation in which $M = 5$ min, $T = 1$ min, and $S = \$300/\text{hr}$, the operators choose $N = 50$. Assuming that this is an optimal choice, find the unit cost of the user's waiting time W .

2.3 In the Exec II batch system, users would submit a large number of jobs in the morning. These jobs took hours to complete and thereby prevented fast response. Suggest a modification of the scheduling policy that would discourage users from doing this.

2.4 In the Scope system for the CDC 6600 computer, system resources (processor time, storage, etc.) can remain idle while running jobs wait for operators to mount magnetic tapes. Suggest a solution to this problem.

2.5 Measurements on the CTSS system showed that about half of all user requests could be classified as file manipulation, program input, and editing. How would you use this information about expected workload to improve processor utilization at a reasonable cost without degrading user response?

CHAPTER 3 PROCESS DESCRIPTION AND CONTROL

- 3.1** Name five major activities of an OS with respect to process management, and briefly describe why each is required.
- 3.2** Consider a computer with N processors in a multiprocessor configuration.
- a.** How many processes can be in each of the Ready, Running, and Blocked states at one time?
 - b.** What is the minimum number of processes that can be in each of the Ready, Running, and Blocked states at one time?
- 3.3** Consider a system in which the following states are defined for processes: Execute (running), Active (ready), Blocked, and Suspend. A process is blocked if it is waiting for permission to use a resource, and it is suspended if it is waiting for an operation to be completed on a resource it has already acquired. In many operating systems, these two states are lumped together as the blocked state, and the suspended state has the definition we have used in this chapter. Compare the relative merits of the two sets of definitions.
- 3.4** What would happen if you executed the following piece of code:
- ```
main()
{
 for(;;)
 fork();
}
```
- Note:* Don't try this on a real machine.
- 3.5** During a process switch, the operating system executes instructions that choose the next process to execute. These instructions are typically at a fixed location in memory. Why?

# CHAPTER 5 CONCURRENCY: MUTUAL EXCLUSION AND SYNCHRONIZATION

- 5.1** In a multiprocessing environment, a single dispatch list may be accessed by each multiprocessor's dispatcher. Give an example of a race condition that can occur if two dispatchers access the dispatch list concurrently.
- 5.2** Nicole Kidman, Penelope Cruz, and Mimi Rogers, like all members of Hollywood, often eat at Spago's. But Nicole, Penelope, and Mimi would rather NOT eat at Spago's on the same evening as one of the other two. Think of Spago's as a system resource (e.g., file). Think of Nicole, Penelope, and Mimi as three processes needing access to this resource. Does the following solution to the critical section problem for three processes satisfy the desired mutual exclusion requirements (mutual exclusion, no deadlock, no starvation, process can enter empty critical section without delay)? Justify your response by discussing each of the critical section requirements, if the solution is correct; or one of the critical section requirements, if the solution is not correct.

have three flags (shared variables): Cole, Lope, and Mimi  
 have neon sign (shared variable) which flashes either  
 "Nicole" or "Penelope" or "Mimi-R"  
 initially: all flags are down and sign flashes "Nicole"

| Nicole                                                                                                                                                                                               | Penelope                                                                                                                                                                                         | Mimi                                                                                                                                                                                                    |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <pre>repeat   [play with kids];   raise Cole;   SIGN = Nicole;   while (((Lope is up)     OR (Mimi is up))     AND (SIGN == Nicole))      [do nothing];   [eat dinner];   lower Cole; forever;</pre> | <pre>repeat   [call Tom];   raise Lope;   SIGN = Penelope;   while (((Cole is up)     OR (Mimi is up))     AND (SIGN == Penelope))    [do nothing];   [eat dinner];   lower Lope; forever;</pre> | <pre>repeat   [make a movie];   raise Mimi;   SIGN = Mimi-R;   while (((Cole is up)     OR (Lope is up))     AND (SIGN == Mimi- R))          [do nothing];   [eat dinner];   lower Mimi; forever;</pre> |



- 5.3** Does the following solution to the critical section problem for two processes satisfy the desired mutual exclusion requirements (mutual exclusion, no deadlock, no starvation, process can enter empty critical section without delay)? Justify your response by discussing each of the critical section requirements, if the solution is correct; or one of the critical section requirements, if the solution is not correct.
- have two flags (shared variables): Cole and Lope  
 have neon sign (shared variable) which flashes either "Nicole" or "Penelope"  
 initially: both flags are down and sign flashes "Nicole"

| Nicole                                                                                                                                                                                                                                                             | Penelope                                                                                                                                                                                                                                            |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <pre>repeat   [play with kids];   raise Cole;   while (Lope is up) {     if (SIGN == Penelope) {       lower Cole;       while (SIGN == Penelope)         [do nothing];       raise Cole;     }   }   [eat dinner];   SIGN = Penelope   lower Cole; forever;</pre> | <pre>eat   [call Tom];   raise Lope;   while (Cole is up) {     if (SIGN == Nicole) {       lower Lope;       while (SIGN == Nicole)         [do nothing];       raise Lope;     }   }   [eat dinner];   SIGN = Nicole   lower Lope; forever;</pre> |

- 5.4** Some do not agree with the name of the semaphore's wait operation. They do not like the word: *wait*. Explain.
- 5.5** Consider the following functions that are shared by multiple processes:

```
static int count = 0;
int increment(void)
{
 count++;
 if (count > 5) {
 printf("counter %d reached value > 5", count);
 return 0;
 }
 return 1;
}

int decrement(void)
{
 while (count > 5) {
 printf("counter %d is > 5:", count);
 count --;
 }
 if (count == 0) return 0;
 else return 1;
}
```

}

Use any synchronization primitives (semaphores, mutex) to make the two functions atomic.

- 5.6** Another atomic machine instruction that supports mutual exclusion that is often mentioned in the literature is the `test&set` instruction, defined as follows:

```
boolean test_and_set (int i)
{
 if (i == 0) {
 i = 1;
 return true;
 }
 else return false;
```

Define a procedure similar to those of Figure 5.2 that uses the `test&set` instruction.

- 5.7** In the discussion of the producer/consumer problem with finite buffer (Figure 5.12), note that our definition allows at most  $n - 1$  entries in the buffer.
- a. Why is this?
  - b. Modify the algorithm to remedy this deficiency.

- 5.8** The Busy Bank Problem. A bank has  $n$  tellers, all serving customers at teller windows.  $k$  of the  $n$  tellers are designated as "quick service" tellers, who only process single withdrawal transactions. Customers arriving at the bank join one of two queues, the "normal service" queue, or the "quick service" queue. They may only join the latter queue if they wish to make a single withdrawal transaction. Customers in the quick service queue are processed by the  $k$  quick service tellers, unless the normal service queue is empty, when they can also be processed by any of the normal service tellers. Customers in the normal service queue are processed by the  $(n-k)$  normal service tellers, unless the quick service queue is empty, when they can also be processed by any of the  $k$  quick service tellers.

Write a suite of suitable synchronized concurrent processes to model the processing of customers by the bank. You may use any (formal) programming language notation you prefer, as long as you clearly identify it. *Hint:* study the structure of the barbershop problem in Appendix A.

- 5.9** Consider a concurrent program with two processes, p and q, defined as follows. A, B, C, D, and E are arbitrary atomic (indivisible) statements. Assume that the main program (not shown) does a **parbegin** of the two processes.

```
void p()
{
 A;
 B;
 C;
}

void q()
{
 D;
 E;
}
```

Show all the possible interleavings of the execution of the preceding two processes (show this by giving execution "traces" in terms of the atomic statements).

- 5.10** Show where/how to add semaphores to the program in the preceding problem to insure that the printf() is never executed. Your solution should allow as much concurrency as possible.
- 5.11** Suppose that an operating system does not have access to an atomic test-and-set, compare and swap, or a lock instruction, but does have the ability to mask out interrupts with a spl(x) routine which sets the processor priority level to x; the return value of spl() is the old processor priority level. spl(x) is a privileged instruction (can only be executed in kernel mode) and is atomic. Sketch how you would make/implement an operating system service (system call) for the user that would allow her/him to easily implement mutually exclusive access to a critical section. A busy-waiting solution to the problem is OK, but remember that busy waiting inside the kernel will hang the system.
- 5.12** In the discussion of Figure 5.9, the book makes the following comment: " It would not do simply to move the conditional statement inside the critical section of the consumer because this could lead to deadlock (e.g., after line 8 of the table)." Explain how this could lead to deadlock.

# CHAPTER 6 CONCURRENCY: DEADLOCK AND STARVATION

**6.1** Consider the following snapshot of a system. There are no outstanding unsatisfied requests for resources.

|           |    |    |    |  |
|-----------|----|----|----|--|
| available |    |    |    |  |
| r1        | r2 | r3 | r4 |  |
| 1         | 5  | 2  | 0  |  |

|         |                    |    |    |    |                |    |    |    |
|---------|--------------------|----|----|----|----------------|----|----|----|
|         | current allocation |    |    |    | maximum demand |    |    |    |
| process | r1                 | r2 | r3 | r4 | r1             | r2 | r3 | r4 |
| p0      | 0                  | 0  | 1  | 2  | 0              | 0  | 1  | 2  |
| p1      | 1                  | 0  | 0  | 1  | 1              | 5  | 1  | 1  |
| p2      | 1                  | 3  | 0  | 4  | 2              | 3  | 5  | 6  |
| p3      | 0                  | 6  | 3  | 2  | 0              | 6  | 7  | 2  |
| p4      | 0                  | 0  | 1  | 4  | 0              | 6  | 5  | 6  |

Is the system safe?

**6.2** Apply the deadlock detection algorithm of Section 6.4 to the following data and show the results.

$$\text{Available} = (2 \ 1 \ 0 \ 0)$$

$$\text{Request} = \begin{pmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{pmatrix} \quad \text{Allocation} = \begin{pmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{pmatrix}$$

- 6.3** In the context of Dijkstra's Banker's Algorithm, discuss whether the following state is safe or unsafe. If safe, show how it is possible that all processes complete.

| available |  |    |    |  |
|-----------|--|----|----|--|
|           |  | A  | B  |  |
|           |  | 10 | 15 |  |

| proces<br>s | current<br>allocation |   | maximum<br>demand |   |
|-------------|-----------------------|---|-------------------|---|
|             | A                     | B | A                 | B |
| User 1      | 2                     | 3 | 10                | 5 |
| User 2      | 3                     | 3 | 8                 | 5 |
| User 3      | 2                     | 2 | 4                 | 4 |
| User 4      | 2                     | 5 | 4                 | 8 |

- 6.4** Consider the following snapshot of a system. There are no outstanding unsatisfied requests for resources.

| available |    |    |    |  |
|-----------|----|----|----|--|
| r1        | r2 | r3 | r4 |  |
| 2         | 1  | 0  | 0  |  |

| process | current<br>allocation |    |    |    | maximum<br>demand |    |    |    | still needs |    |    |    |
|---------|-----------------------|----|----|----|-------------------|----|----|----|-------------|----|----|----|
|         | r1                    | r2 | r3 | r4 | r1                | r2 | r3 | r4 | r1          | r2 | r3 | r4 |
| p1      | 0                     | 0  | 1  | 2  | 0                 | 0  | 1  | 2  |             |    |    |    |
| p2      | 2                     | 0  | 0  | 0  | 2                 | 7  | 5  | 0  |             |    |    |    |
| p3      | 0                     | 0  | 3  | 4  | 6                 | 6  | 5  | 6  |             |    |    |    |
| p4      | 2                     | 3  | 5  | 4  | 4                 | 3  | 5  | 6  |             |    |    |    |
| p5      | 0                     | 3  | 3  | 2  | 0                 | 6  | 5  | 2  |             |    |    |    |

- Compute what each process still might request and display in the columns labeled "still needs."
- Is this system currently in a safe or unsafe state? Why?
- Is this system currently deadlocked? Why or why not?
- Which processes, if any, are or may become deadlocked?
- If a request from p3 arrives for (0, 1, 0, 0), can that request be safely granted immediately? In what state (deadlocked, safe, unsafe) would immediately granting that whole request leave the system? Which processes, if any, are or may become deadlocked if this whole request is granted immediately?

**6.5** For this problem, consider the following ways of handling deadlock.

- (i) banker's algorithm
- (ii) detect deadlock and kill thread, releasing all its resources
- (iii) reserve all resources in advance
- (iv) restart thread and release all resources if thread needs to wait
- (v) resource ordering
- (vi) detect deadlock and roll back thread's actions

- a.** One criteria to use in evaluating different approaches to deadlock is which permits the greatest concurrency -- in other words, which allows the most threads to make progress without waiting when there isn't deadlock. Give a rank order from 1 to 6 for each of the ways of handling deadlock listed below, where 1 allows the greatest degree of concurrency, and 6 allows the least concurrency; if two approaches offer roughly equal concurrency, indicate that.
- b.** Another criteria to use in evaluating deadlock algorithms is efficiency -- in other words, which requires the least CPU overhead. Rank order the approaches according to efficiency, with 1 being the most efficient, assuming that encountering deadlock is a very rare event. Again, indicate if two approaches are equally efficient. Does your ordering change if deadlocks happen frequently?

**6.6 a.** Three processes share four resource units that can be reserved and released only one at a time. Each process needs a maximum of two units. Show that a deadlock cannot occur.

- b.**  $N$  processes share  $M$  resource units that can be reserved and released only one at a time. The maximum need of each process does not exceed  $M$ , and the sum of all maximum needs is less than  $M + N$ . Show that a deadlock cannot occur.

## CHAPTER 7 MEMORY MANAGEMENT

- 7.1** Assume that we are using OS/MVT (variable size partitions) and we have batch processes arriving in order, with the following requests for storage:

20k, 30k, 10k, 100k, 60k

Assume that we have four holes in memory of size: 50k, 30k, 200k, 16k, 30k. Where would each of the First-fit, Best-fit and Worst-fit algorithms place the processes?

- 7.2** Consider the following segment table:

| Segment | Base | Length |
|---------|------|--------|
| 0       | 1219 | 600    |
| 1       | 3300 | 14     |
| 2       | 90   | 100    |
| 3       | 2327 | 580    |
| 4       | 1952 | 96     |

What are the physical addresses for the following logical addresses?

- a. 0, 4302    b. 1, 15    c. 2, 50    d. 3, 400    e. 4, 112

- 7.3** Suppose a fixed partitioning memory system with partitions of 100K, 500K, 200K, 300K, and 600K (in memory order) exists. All five partitions are currently available.
- Using the best fit algorithm, show the state of memory after processes of 212K, 417K, 112K, and 350K (in request order) arrive.
  - Using the best available algorithm, show the state of memory after processes of 212K, 417K, 112K, and 350K (in request order) arrive.
  - At the end of part (a), how much internal fragmentation exists in this system?
  - At the end of part (b), how much external fragmentation exists in this system?

- 7.4** Consider a logical address space of eight pages of 1024 bytes each, mapped onto a physical memory of 32 frames.
- How many bits are there in the logical address?
  - How many bits are there in the physical address?
- 7.5** A system uses (contiguous) dynamic partition memory management, with 110K of memory for user processes. The current memory allocation table is shown below:

| Job | Base Address | Length |
|-----|--------------|--------|
| A   | 20           | 10     |
| B   | 46           | 18     |
| C   | 90           | 20     |

A new job, D, arrives needing 15K of memory. Show the memory allocation table entry for job D for each of the following memory allocation strategies: first fit, best fit, worst fit.

- 7.6** Consider a system with a 16KB memory. The sequence of processes loaded in and leaving the memory are given in the following.

P1 7K loaded  
P2 4K loaded  
P1 terminated and returned the memory space  
P3 3K loaded  
P4 6K loaded


Assume that when a process is loaded to a selected "hole", it always starts from the smallest address. E.g. P1 will be loaded in memory locations 0 to 8K – 1 since the entire memory is one hole. Give the memory map showing allocated portion and free portion after the end of the sequence (if a process cannot be loaded, indicate that) for the following placement algorithms. Also, indicate the internal/external fragmentations.

- first fit
  - best fit
  - buddy
  - simple paging (assume that each page is of size 2K)
- 7.7** On a system with  $2^{32}$  bytes of memory and fixed partitions with a partition size of  $2^{20}$  bytes, what is the minimum number of bits needed in an entry in the process table to record the partition to which a process has been allocated?



**7.8** On a system using fixed partitions with partition sizes of  $2^8$ ,  $2^{24}$ , and  $2^{64}$  bytes, how many bits must the bounds register (Figure 7.8) have?

**7.9** A dynamic partitioning scheme is being used, and the following is the memory configuration at a given point in time:



The shaded areas are allocated blocks; the white areas are free blocks. The next three memory requests are for 40M, 20M, and 10M. Indicate the starting address for each of the three blocks using the following placement algorithms:

- First-fit
- Best-fit
- Next-fit. Assume the most recently added block is at the beginning of memory.
- Worst-fit

**7.10** A virtual address  $a$  in a paging system is equivalent to a pair  $(p, w)$ , in which  $p$  is a page number and  $w$  is a byte number within the page. Let  $z$  be the number of bytes in a page. Find algebraic equations that show  $p$  and  $w$  as functions of  $z$  and  $a$ .

## CHAPTER 8 VIRTUAL MEMORY

**8.1** Consider a demand paging system with the following time-measured utilizations.

processor utilization 20 %

paging disk 97.7 %

other I/O devices 5 %

Which of the following if any will probably improve processor utilization

- a. Install a faster processor.
- b. Install a bigger paging disk
- c. Increase the degree of multiprogramming
- d. Install more main memory
- e. Decrease the time quantum allocated to each process

**8.2** What is the probable cause of thrashing?

- a. I/O drivers that are not properly connected to the I/O subsystem.
- b. Repeated parity errors
- c. Disk crashes
- d. A local page replacement algorithm
- e. Process (e.g., being unable to establish their working set of pages
- f. A FIFO page replacement algorithm

**8.3** Consider a paging system with a page table stored in memory.

- a. If a memory reference takes 200 nanoseconds, how long does a paged reference take (reference is retrieved from page table in memory and then program data is retrieved from memory) Assume all pages are in memory?
- b. If we add associative registers (TLB), and 75 percent of all page-table references are found in the associative registers, what is the effective memory reference time on average? (Assume that finding a page table entry in the associative registers takes 10 nanoseconds, if the entry is there.) Assume all pages are in memory.
- c. What is the effective memory reference time if the hit ratio for the associative registers is 70%, the hit ratio for memory is 20%, and 10% of the time we must go out to the disk to load the faulted page in memory, where context switching, disk transfer time, etc takes 100 m seconds.).

**8.4** Assume a demand paged memory management system with the following characteristics:

page size = 4K =  $2^{12}$  bytes

physical address space =  $2^{24}$  bytes

logical address space =  $2^{32}$  bytes

TLB size =  $2^6$  bytes

- a. How many bits are needed to specify a logical address? Give the number of bits for the page number and offset.
- b. How many page table entries can fit in the TLB if each entry only contains the information needed for logical to physical translation. Show your work.
- c. How many page table entries are needed when a program uses its full logical address space?
- d. Part (c) indicates a serious problem that arises from having a very large logical address space. What is this problem and how could an OS solve it? Discuss the consequences of your solution for runtime overhead during program execution (just a few sentences).

**8.5** Suppose:

TLB lookup time = 20 ns

TLB hit ratio = 80%

memory access time = 75 ns

swap page time = 500,000 ns

50% of pages are dirty

OS uses a single level page table

- a. What is the effective access time (EAT) if we assume the page fault rate is 0%? (Show your work.) Assume the cost to update the TLB, the page table, and the frame table (if needed) is negligible.
- b. What is the effective access time (EAT) if we assume the page fault rate is 10%? (Show your work.) Assume the cost to update the TLB, the page table, and the frame table (if needed) is negligible.

- 8.6** A process has four page frames allocated to it. (All the following numbers are decimal, and everything is numbered starting from zero). The time of the last loading of a page into each page frame, the time of last access to the page in each page frame, the virtual page number in each page frame, and the referenced (R) and modified (M) bits for each page frame are as shown (the times are in clock ticks from the process start at time 0 to the event — not the number of ticks since the event to the present).

| Virtual page number | Page frame | Time loaded | Time referenced | R bit | M bit |
|---------------------|------------|-------------|-----------------|-------|-------|
| 2                   | 0          | 60          | 161             | 0     | 1     |
| 1                   | 1          | 130         | 160             | 1     | 0     |
| 0                   | 2          | 26          | 162             | 1     | 0     |
| 3                   | 3          | 20          | 163             | 1     | 1     |

A page fault to virtual page 4 has occurred at time 164. Which page frame will have its contents replaced for each of the following memory management policies? Explain why in each case.

- FIFO (first-in-first-out)
- LRU (least recently used)
- Clock
- Optimal (Use the following reference string.)
- Given the aforementioned state of memory just before the page fault, consider the following virtual page reference string:

4, 0, 0, 0, 2, 4, 2, 1, 0, 3, 2

How many page faults would occur if the working set policy with LRU were used with a window size of 4 instead of a fixed allocation? Show clearly when each page fault would occur.

- 8.7** Consider the following page reference string:

0, 1, 7, 0, 1, 2, 0, 1, 2, 3, 2, 7, 1, 0, 3, 1, 0, 3

How many page faults would occur for the following replacement algorithms? Assume no prepaging occurs, and show what pages are in memory at each given time.

- OPT replacement (three frames are allocated to the process)
- FIFO replacement (three frames are allocated to the process)
- Pure LRU replacement (three frames are allocated to the process)
- Clock Policy (three frames are allocated to the process)

**8.8** Suppose an instruction takes 1 nanosecond to execute (on average), a page fault takes 20 microseconds of processor time, and it takes 300 microseconds of disk time to read or write a single page. Suppose that on average, 1/2 of the pages are modified. What is the average number of instructions between page faults that would cause the disk to be busy doing page transfers all the time?

**8.9** A virtual memory system exhibits the follow trace of page numbers:

1 2 3 2 6 3 4 1 5 6 1 6 4 2

Simulate the page replacements for the following scenarios:

- a. FIFO with 3 page frames
- b. FIFO with 4 page frames
- c. LRU with 3 page frames
- d. LRU with 4 page frames

and show the number of page faults that occur with each. Does Belady's Anomaly occur? Belady's anomaly states that it is possible to have more page faults when increasing the number of page frames under certain conditions. This was reported in "An Anomaly in Space-Time Characteristics of Certain Programs Running in a Paging Machine," by Belady et al, *Communications of the ACM*, June 1969.

**8.10** Consider a paged virtual memory system with 32-bit virtual addresses and 1K-byte pages. Each page table entry requires 32 bits. It is desired to limit the page table size to one page.

- a. How many levels of page tables are required?
- b. What is the size of the page table at each level? *Hint:* One page table size is smaller.
- c. The smaller page size could be used at the top level or the bottom level of the page table hierarchy. Which strategy consumes the least number of pages?

**8.11** Certain algorithms that work very well when all of the data fit in memory do not work so well when the virtual memory exceeds the available physical memory, due to the amount of time that may be spent on page replacement. The following program multiplies two matrices, where each matrix is stored as a two dimensional array. The array is stored in row-order, meaning that consecutive elements in the same row of the array will be stored in consecutive locations in memory.

```
multiply(A,B,C)
{
 int A[SIZE][SIZE]. B[SIZE][SIZE], C[SIZE][SIZE];
 int i,j,k,temp;

 for all i /* this done in parallel */
 for (j=0; j<SIZE; j++) {
 temp = 0;
 for (k=0; k<SIZE; k++)
 temp += a[i][k] * b[k][j];
 c[i][j] = temp;
 }
}
```

Suppose the above program was written for a shared-memory multiprocessor. Each processor has a number assigned to it to use for the variable *i*. Therefore, each processor is responsible for computing a row of *C*. Furthermore, assume that they all execute in parallel at the same rate. We will be interested in the performance of this program in light of very large arrays, which will not all fit in memory at the same time. For this problem, assume an optimal page replacement algorithm, where 1000 pages are available for the arrays (you can ignore *i,j,k* and the source program). Each page may hold up to 1000 elements of any array.

Assume there are 100 processors, assigned numbers from 0 to 99. Each will be responsible for doing a set of rows in the multiplication. For example, when multiplying 1000x1000 matrices, Processor 0 will find the first 10 rows, Processor 1 the next 10, and so on. The code each processor runs will be as follows:

```
int A[SIZE][SIZE]. B[SIZE][SIZE], C[SIZE][SIZE];
multiply(A,B,C)
{
 int i,j,k,temp;
 int mult = SIZE / 100;

 for (i=mult*ProcN ; i<mult * (ProcN + 1); i++)
 for (j=0; j<SIZE; j++)
```

```

 {
 temp = 0;
 for (k=0; k<SIZE; k++)
 temp += a[i][k] * b[k][j];
 c[i][j] = temp;
 }
}

```

It is assumed that these are running roughly in parallel, though their memory accesses will be serialized by the memory unit. Therefore, accesses to the arrays will look something like this:

```

{ a[0][0], a[10][0], a[20][0], a[990][0] } one access
{ b[0][0], b[0][0], b[0][0], b[0][0] } for each
{ a[0][1], a[10][1], a[20][1], a[990][1] } processor
{ b[1][0], b[1][0], b[1][0], b[1][0] }
etc.

```

An Optimal Page Replacement Policy is one that accomplishes all this in the fewest page faults -- based on clairvoyantly predicting what will be needed soon, and should not be replaced. Here, for example, several pages of A are used repeatedly, so should hang about in memory for a long time.

The circumstances to consider are the following:

- a.** SIZE = 1,000 (total memory required is 3 times that available)
  - b.** SIZE = 10,000 (total memory required 300 times that available)
- Estimate how many page faults occur in each case, (one significant digit is enough) and identify where most of them come from.

**8.12** For this problem, assume a demand paged virtual memory system with a page size of 100 words (decimal values throughout). A process running on this system generates a sequence of logical addresses, given in the table below.

|    |    |     |     |    |     |     |     |     |     |     |     |
|----|----|-----|-----|----|-----|-----|-----|-----|-----|-----|-----|
| 10 | 11 | 104 | 170 | 73 | 309 | 185 | 245 | 246 | 434 | 458 | 364 |
|----|----|-----|-----|----|-----|-----|-----|-----|-----|-----|-----|

Assume the process is allotted exactly two page frames, and that none of its pages are resident (in page frames) when the process begins execution.

- a.** Determine the page number corresponding to each logical address and fill them into a table with one row and 12 columns. This is often called a reference string for the process.
- b.** Consider the reference string determined in part (a). Determine which references result in page faults, assuming FIFO page replacement is used, indicating your conclusions by placing Fs in the corresponding cells of a table with one row and 12 columns. The page fault rate is the number of page faults divided by the total

number of references made. What is the page fault rate for this case? Round your answer to an integer percentage.

- c. Repeat (b) for LRU.
- d. Repeat (b) for optimal page replacement.

**8.13** Assume a system that has 224 words of virtual memory and 218 words of physical memory, with a page size of 28 words. All addresses in this system are expressed in hexadecimal. A particular process running on this system is allocated 2700 words of virtual memory. The current page map table for that process is shown below:

| Logical Page # | Frame # | VM Page # |
|----------------|---------|-----------|
| 000            | 212     | 0021      |
| 001            | 002     | 01A0      |
| 002            |         | 7FA3      |
| 003            | 3B2     | 7FA4      |
| 004            | 3F1     | 7FA9      |
| 005            |         | 901B      |
| 006            |         | A113      |
| 007            |         | CA00      |
| 008            |         | CA01      |
| 009            | 012     | CA02      |
| 00A            |         | CA03      |

- a. How many pages of virtual memory does this system have? Give a simplified decimal answer.
- b. How many page frames does this system have? Give a simplified decimal answer.

In each of the following questions you are given a logical address generated by this process. Determine the virtual and physical addresses that correspond to each logical address. If there is no valid virtual or physical address indicate why. Express both physical and virtual addresses in hexadecimal.

- c. 3F1
- d. 214
- e. A90



## CHAPTER 9 UNIPROCESSOR SCHEDULING

**9.1** Consider the following set of processes:

| Process Name | Arrival Time | Processing Time |
|--------------|--------------|-----------------|
| A            | 0            | 1               |
| B            | 1            | 9               |
| C            | 2            | 1               |
| D            | 3            | 9               |

Perform the same analysis as depicted in Table 9.5 and Figure 9.5 for this set.

**9.2** Suppose the following jobs are to be executed in a uniprocessor system.

| Job Number | Arrival Time | Service Time |
|------------|--------------|--------------|
| 1          | 0            | 4            |
| 2          | 1            | 8            |
| 3          | 3            | 2            |
| 4          | 10           | 6            |
| 5          | 12           | 5            |

Assume the overhead of context switching is one time unit. For each of the following scheduling methods, give (i) a timing chart to illustrate the execution sequence, (ii) the average job turnaround time, (iii) the normalized turnaround time for each job, and (iv) the processor efficiency.

**a.** FCFS

**b.** SPN

**c.** SRTN

**d.** RR, quantum = 3

**e.** Multilevel Feedback Queue with queues numbered 1-10, quantum =  $2^i$ , where  $i$  is the queue level number and processes are initially placed in the first queue (i.e., level 1). In this scheduling policy, each process executes at a particular level for one quantum and then moves down a level; processes never move up a level.

- 9.3** Most round-robin schedulers use a fixed size quantum. Give an argument in favor of a small quantum. Now give an argument in favor of a large quantum. Compare and contrast the types of systems and jobs to which the arguments apply. Are there any for which both are reasonable?
- 9.4** Consider a page reference string for a process with a working set of  $m$  frames, initially all empty. The page reference string is of length  $p$  with  $n$  distinct page numbers in it. For any page replacement algorithm:
- a.** What is a lower bound on the number of page faults?
  - b.** What is an upper bound on the number of page faults?

# CHAPTER 11 I/O MANAGEMENT AND DISK SCHEDULING

**11.1** For the following disk accesses, compute the number of head movements for the following list of seeks to disk cylinder: 26 37 100 14 88 33 99 12

Assume head is initially positioned over 26.

- a. FCFS      b. SSTF      c. SCAN (going up)      d. C-SCAN (going up)

**11.2** Although DMA does not use the CPU, the maximum transfer rate is still limited. Consider reading a block from the disk. Name three factors that might ultimately limit the rate transfer.

**11.3** Let us assume a disk with rotational speed of 15,000 rpm, 512 bytes per sector, 400 sectors per track and 1000 tracks on the disk, average seek time is 4ms. We want to transmit a file of size 1 MByte, which is stored contiguously on the disk.

- a. What is the transfer time for this file?
- b. What is the average access time for this file?
- c. What is the rotational delay in this case?
- d. What is the total time to read 1 sector?
- e. What is the total time to read 1 track?

**11.4 Spooling** is an I/O technique developed in the early days of multiprogramming systems. It is a way of dealing with a dedicated I/O device. Each user process generates an I/O stream for a device that is directed to a disk file rather than to the device itself. The responsibility for moving data between the disk and the required device is delegated to a separate process, called a spooler. Spooling is appropriate for devices that cannot accept interleaved data streams, such as a printer.

**Prefetching** is a method of overlapping the I/O of a job with that job's own computation. The idea is simple. After a read operation completes and the job is about to start operating on the data, the input device is instructed to begin the next read immediately. The processor and input device are then both busy. With luck, by the time that the job is ready for the next data item, the input device will have finished reading that data item. The processor can then begin processing the newly read data, while the input device starts to read the following data.

A similar idea can be used for output. In this case, the job creates data into a buffer until an output device can accept them.

Compare the prefetching scheme with the spooling scheme, where the processor overlaps the input of one job with the computation and output of other jobs.

- 11.5** Consider a disk system with 8 sectors per track and 512 bytes per sector. The disk rotates at 3000 rpm and has an average seek time of 15 msec. Also, consider a file consisting of 8 blocks. Compute the total time for accessing the entire file if the following allocation algorithms are used.
- a.** contiguous allocation,
  - b.** indexed allocation
- 11.6** What are the advantages and disadvantages of smaller vs. larger disk sector size?
- 11.7** Most disks today use a fixed sector size. Many IBM mainframe disks allow the programmer to determine the block (sector) size. Explain advantages and disadvantages of allowing a variable sector size.
- 11.8** Let us assume a disk with rotational speed of 15,000 rpm, 512 bytes per sector, 400 sectors per track and 1000 tracks on the disk, average seek time is 4ms. We want to transmit a file of size 1 MByte, which is stored contiguously on the disk.
- a.** What is the transfer time for this file?
  - b.** What is the average access time for this file?
  - c.** What is the rotational delay in this case?
  - d.** What is the total time to read 1 sector?
  - e.** What is the total time to read 1 track?
- 11.9** A magnetic disk with 5 platters has 2048 tracks/platter, 1024 sectors/track (fixed number of sectors per track), and 512-byte sectors. What is its total capacity?

## CHAPTER 12 FILE MANAGEMENT

- 12.1** A UNIX i-node has 13 direct pointers, 1 indirect pointer, 1 double indirect pointer and 1 triple indirect pointer. Assume that each 32-bit pointer identifies one block of 8KB. How large a file can the i-node handle?
- 12.2** A sequential access file has fixed-size 150-byte records. Assuming the first record is record 1, the first byte of record 5 will be at what logical location?
- 12.3** A direct access file has fixed-size 150-byte records. Assuming the first record is record 1, the first byte of record 5 will be at what logical location?
- 12.4** On a system using 60-byte records and 500-byte blocks, how much space will be wasted in each block?
- 12.5** Consider a system using unspanned blocking and 100-byte blocks and a file containing records of 20, 50, 35, 70, 40, 20 bytes. What percentage of space will be wasted in the blocks allocated for the file?